
All-Pairs Shortest Paths in Spark

Charles Y. Zheng and Jingshu Wang

Department of Statistics
Stanford University
Stanford, CA 94305

{snarles, jinshuw}@stanford.edu

Arzav Jain

Department of Computer Science
Stanford University
Stanford, CA 94305

arzavj@cs.stanford.edu

1 Path finding in distributed APSP

The output of the original algorithm is a matrix of shortest distances $S \in \mathbb{R}^{n \times n}$ where each S_{ij} is the shortest distance from node i to j . Here we want to add a path lookup function $\text{FindPath}(i, j)$ which returns for a pair of node (i, j) the shortest path itself from node i to node j .

There are three approaches to consider:

- Calculate $\text{FindPath}(i, j)$ directly from the distance matrix S
- Store one midpoint for each (i, j) pair as a matrix $M \in \mathbb{R}^{n \times n}$ in the distributed block APSP algorithm, and then calculate $\text{FindPath}(i, j)$ from M
- For each (i, j) pair, store two midpoints m_1, m_2 in a three-dimensional array $M \in \mathbb{R}^{n \times n \times 2}$ in the distributed block APSP algorithm, and then calculate $\text{FindPath}(i, j)$ from M .

Both the first two approaches need n iterations in calculating $\text{path}(i, j)$ for the worst case, while we will show in this session that the third approach guarantees the number of iterations to be at most $\lceil \log_2 n \rceil$ with properly chosen midpoints.

1.1 Criteria for choosing the midpoints

For an (i, j) pair with its shortest path $i \rightarrow k_1 \rightarrow k_2 \rightarrow \dots \rightarrow k_{L-1} \rightarrow j$, define its path length as $l_{ij} = L$. We require the midpoints $m_1 = M_{ij1}$ and $m_2 = M_{ij2}$ to satisfy

1. $m_1, m_2 \in \{i, k_1, k_2, \dots, k_{L-1}, j\}$
2. $l_{ij} = l_{im_1} + l_{m_1 m_2} + l_{m_2 j}$
3. $\max(l_{im_1}, l_{m_1 m_2}, l_{m_2 j}) \leq \max(l_{ij}/2, 1)$

If M satisfies the above criteria, then the number of iterations in the lookup function $\text{path}(i, j)$ will be at most $\lceil \log_2 n \rceil$. More details can be found in ??.

1.2 Algorithm for updating the midpoints in distributed APSP

The initialization of M is

$$M_{ij1}^{(0)} = M_{ij2}^{(0)} = \begin{cases} i & \text{if } (i \rightarrow j) \in E \text{ or } i = j \\ \star & \text{if } (i \rightarrow j) \notin E \end{cases}$$

where $\star \notin V$ is some symbol to denote an invalid midpoint. To properly update midpoints in our distributed block APSP algorithm, we need to store and update another three-dimensional array $W \in \mathbb{R}^{n \times n \times 3}$ which stores for each (i, j) pair and midpoints (m_1, m_2) the current path lengths l_{im_1} , $l_{m_1 m_2}$ and $l_{m_2 j}$. The initialization of W is

$$W_{ij1}^{(0)} = W_{ij2}^{(0)} = \begin{cases} 0 & \text{if } (i \rightarrow j) \in E \text{ or } i = j \\ \infty & \text{if } (i \rightarrow j) \notin E \end{cases}$$

$$W_{ij3}^{(0)} = \begin{cases} 1 & \text{if } (i \rightarrow j) \in E \text{ or } i = j \\ \infty & \text{if } (i \rightarrow j) \notin E \end{cases}$$

For a path $i \rightarrow \dots \rightarrow j$, denote $v_{ij} = (m_1, m_2, l_{im_1}, l_{m_1m_2}, l_{m_2j})$. Then for joining two paths $i \rightarrow \dots \rightarrow k$ and $k \rightarrow \dots \rightarrow j$, we define the following function $\text{MERGE}(v_{ik}, v_{kj}, k)$ to get v_{ij} for the joint path $i \rightarrow \dots \rightarrow k \rightarrow \dots \rightarrow j$:

Algorithm 1 Merge midpoints of two adjacent paths

function $\text{MERGE}(v_1 = (m_1, m_2, l_1, l_2, l_3), v_2 = (m_4, m_5, l_4, l_5, l_6), m_3)$
 $l = \sum_{i=1}^6 l_i$
for $t = 1, 2, 3, 4$ **do**
 if $\sum_{i=1}^t l_i \leq l/2$ & $\sum_{i=1}^{t+1} l_i \geq l/2$ **then**
 Break
 end if
end for
Return $v = (m_j, m_{t+1}, \sum_{i=1}^t l_i, l_{t+1}, \sum_{i=t+2}^6 l_i)$
end function

We call $v = (m_1, m_2, l_1, l_2, l_3)$ as *flat* if

$$\max(l_1, l_2, l_3) \leq \max(1, (l_1 + l_2 + l_3)/2) < \infty$$

Lemma 1.1. *If v_{ik} and v_{kj} are flat and $i \neq k \neq j$, then $v = \text{MERGE}(v_{ik}, v_{kj}, k)$ is also flat.*

Proof. Let $v_{ik} = (m_1, m_2, l_1, l_2, l_3)$, $v_{kj} = (m_4, m_5, l_4, l_5, l_6)$, $k = m_3$ and $l = \sum_{s=1}^6 l_s$. From $i \neq k \neq j$, we have $l \geq 2$.

As v_{ik} and v_{kj} are *flat*, we have $l_1 \leq \max(1, (l_1 + l_2 + l_3)/2) \leq l/2$ and similarly $l_6 \leq l/2$. Thus, there exists $t \in \{1, 2, 3, 4\}$ that both $\sum_{s=1}^t l_s \leq l/2$ and $\sum_{s=1}^{t+1} l_s \geq l/2$ holds. Also $l_{t+1} \leq \max(1, (l_1 + l_2 + l_3)/2, (l_4 + l_5 + l_6)/2) \leq l/2$, thus $v = \text{MERGE}(v_{ik}, v_{kj}, k)$ is also *flat*. \square

We can now modify the original distributed block APSP algorithm to include updating W and M .

Given an $n \times m$ distance matrix A and an $n \times m$ distance matrix B together with the midpoints matrices (W^A, M^A) and (W^B, M^B) , define a minimum operation as $(C, W^C, M^C) = \min_P((A, W^A, M^A), (B, W^B, M^B))$ by

$$C_{ij} = \min(A_{ij}, B_{ij})$$

$$(M_{ij}^C, W_{ij}^C) = \begin{cases} (M_{ij}^A, W_{ij}^A) & \text{if } C_{ij} = A_{ij} \\ (M_{ij}^B, W_{ij}^B) & \text{if } C_{ij} = B_{ij} \end{cases}$$

Given an $n \times k$ distance matrix A and a $k \times m$ distance matrix B together with the midpoints matrices (W^A, M^A) and (W^B, M^B) , define a *min-plus* product $(C, W^C, M^C) = (A, W^A, M^A) \otimes_P (B, W^B, M^B)$ as

$$C_{ij} = \min_{l=1}^k A_{il} + B_{lj}$$

$$(M_{ij}^C, W_{ij}^C) = \begin{cases} (M_{ij}^A, W_{ij}^A) & \text{if } \arg\min_l (A_{il} + B_{lj}) = j \\ (M_{ij}^B, W_{ij}^B) & \text{if } \arg\min_l (A_{il} + B_{lj}) = i \\ \text{MERGE}((M_{il^*}^A, W_{il^*}^A), (M_{l^*j}^B, W_{l^*j}^B), l^*) & \text{if } l^* = \arg\min_l (A_{il} + B_{lj}) \neq i \text{ or } j \end{cases}$$

for $i = 1, \dots, n$ and $j = 1, \dots, m$.

Also, $\text{APSP}_P(A, M^A, W^A)$ is defined as a modified local APSP method for finding the shortest distance matrix together with the desired midpoints and path lengths matrices.

Here, we give a shorthand description of the modified distributed block APSP including updating W and M , without explicitly specifying the Spark operations.

Algorithm 2 Path-Finding Distributed Block APSP (shorthand)

```
function BLOCKAPSPATH(Adjacency matrix  $A$  given as a BlockMatrix with  $\ell$  row blocks
and  $\ell$  column blocks,  $M^{(0)}, W^{(0)}$ )
   $H^{(0)} \leftarrow (A, M^{(0)}, W^{(0)})$ 
  for  $k = 1, \dots, \ell$  do
    [A-step]
     $H^{kk(k)} \leftarrow \text{APSP}_P(H^{kk(k-1)})$ 
    [B-step]
    for  $i = 1, \dots, \ell, j = 1, \dots, \ell$  do in parallel
      if  $i = k$  and  $j \neq k$  then
         $H^{kj(k)} \leftarrow \min_P(H^{kj(k-1)}, H^{kk(k)} \otimes_P H^{kj(k-1)})$ 
      end if
      if  $i \neq k$  and  $j = k$  then
         $H^{ik(k)} \leftarrow \min_P(H^{ik(k-1)}, H^{ik(k-1)} \otimes_P H^{kk(k)})$ 
      end if
    end for
    [C-step]
    for  $i = 1, \dots, \ell, j = 1, \dots, \ell$  do in parallel
      if  $i \neq k$  and  $j \neq k$  then
         $H^{ij(k)} \leftarrow \min_P(H^{ij(k-1)}, H^{ik(k)} \otimes_P H^{kj(k)})$ 
      end if
    end for
    [D-step]
    if  $k \equiv 0 \pmod q$  then
      Checkpoint  $H^{(k)}$ 
    end if
  end for
  Return  $(S, M, W) = H^{(\ell)}$ , the APSP result tuple
end function
```

1.3 The path lookup function

After obtaining the the midpoints three-dimensional array M , we can efficiently lookup the shortest path of an (i, j) pair of nodes. The lookup function returns a vector of all the other nodes in the path in order except for the starting node i . Note that if there are multiple shortest paths, the algorithm is only able to find one of them.

Algorithm 3 Lookup the path from one node to another

```
function FINDPATH( $i, j$ )
  if  $i == j$  then
    Return NULL
  end if
  if  $M_{ij1} == M_{ij2}$  then
    Return  $j$ 
  end if
  Return (FindPath( $i, M_{ij1}$ ), FindPath( $M_{ij1}, M_{ij2}$ ), FindPath( $M_{ij2}, j$ ))
end function
```

As $\max(l_{iM_{ij1}}, l_{M_{ij1}M_{ij2}}, l_{M_{ij2}j}) \leq \max(1, l_{ij}/2)$, the recursion depth of the above algorithm is upper bounded by $\lceil \log_2 l_{ij} \rceil$, which is at most $\lceil \log_2 n \rceil$ for any node pair in the graph.